

When Can an Agent Stop Remembering? A Pre-Registered Mechanism Study of Context Retirement, Its Audit, and the Cost Regime That Pays

This artifact candidacy draft accompanies the `future-code` research repository. Author names and affiliations require c

September 24, 2026 | Version 7.0 (factorial retirement revision)

Abstract

When can an agent stop remembering its execution history without losing the current evidence and obligations it needs to finish correctly—and in which cost regime does the answer pay? We answer at the mechanism level with a pre-registered, factorial, frozen-protocol study of context retirement: 1,176 runs crossing seven retention policies with a scheduler factor, plus a 160-episode long-horizon sub-cohort, across four registry sizes, three dependency families, and seven seeds, with temporal persistence, obligation lifetimes, three-phase accounting, a five-episode capsule audit against stream ground truth, and two frozen cost regimes (bytes; bytes plus a per-round-trip constant). The result that survives is a *conversion rule with a regime-conditional payoff*. An agent may retire its history only after converting it into verified interface bindings, a typed pending-obligation set, and declared dependencies; lossy conversion measurably fails (periodic summary/reset accepts 28 of 84 runs at the new scale, with 109,761 stale reads and 385 lost obligations), while verified-frontier retirement accepts 84/84 and passes the capsule audit with zero defective service. The payoff is regime-conditional and was pre-registered as such: in the bandwidth regime frontier retirement saves only -8.0% on submitted bytes—no byte-cost win is claimed—but in the round-trip regime, where each on-demand retrieval costs latency-equivalent bytes, it reduces cost per accepted run by 29.6% versus strong observation masking (sensitivity 17.8% to 38.8% across round-trip constants), cutting retrieval operations 44.0% pooled and 45.1% at four times the horizon. The scheduler factor is byte-neutral by design and measured at 0.000% leak with an 87.5% makespan reduction. Grouped certificate layouts survive only as a conditional component selected by measured update density (dense stratum pass; 0.0% decision-rule regret), and a demand-driven adaptive-coordination variant was evaluated against pre-registered promotion gates and *rejected* (-0.0% versus frontier; the correctness-preserving mechanism does not need it). All experiments are synthetic and mechanism-level: no hosted model was called, no token or billing claim is made, and the live-agent cohort remains future work with its protocol pre-registered here.

Keywords: bounded evidence; verification contracts; bandwidth-limited logging; completion hyperedges; cost-aware layout selection.

1 Productivity Boundary and Research Scope

The governing research objective is to reduce the cost and time required for agents to complete large software projects correctly, without increasing coordination failures or weakening acceptance contracts. That objective outruns what any single paper can measure, and this manuscript is explicit about the ladder of evidence it occupies.

What this paper is. It is a mechanism study of context retirement (Section 4) plus the verification and evidence components the surviving rule depends on (Sections 5–11). Two protocol generations were run; the current one is factorial, pre-registered, and git-frozen before data

generation. Five closure rules were declared before data; two of the paper’s own proposed extensions failed their gates and are reported as rejected: the byte-regime productivity claim for frontier retirement (it saves -8.0% on submitted bytes versus masking—no win claimed) and the demand-driven adaptive-coordination variant (-0.0% regime-O versus plain frontier; promotion gate $\geq 10\%$). What survives is (i) a necessary conversion rule—retire history only after converting it into verified bindings, typed obligations, and declared dependencies; lossy conversion fails the acceptance contract at every scale tested (28/84 at the largest grid); (ii) an operations result at equal correctness—frontier retirement cuts on-demand retrieval operations 44.0% and round-trip-regime cost 29.6% versus observation masking, persisting at $4\times$ horizon, with a clean capsule audit; and (iii) a density-selected conditional role for certificate-tree evidence layouts (0.0% decision-rule regret). HACT is not the paper’s center; it survives only where its ablation says it wins.

The boundary that remains. The retirement study is mechanism-level. It models information flow between a context policy and a seeded work stream; it does not call a hosted model, and it cannot establish that any agent finishes faster, bills less, or reasons better. Its staleness, obligation-loss, and audit results are structural predictions about policies, not measurements of agents. The live-agent cohort that could test those predictions has its protocol pre-registered in Section 12.1 and has not been run. The negative controls below—flat checker timings, per-episode verification overhead, unchanged root-status inputs—are reported as limits of the serialization claim, and the retirement cohort’s own closures are reported as limits of the mechanism claim. Nothing in this paper is an agent-productivity measurement, and the abstract says so.

2 Introduction

Long-horizon software agents accumulate execution history, and the accumulation is not free: resubmitting it costs bytes that grow with project length, and naively discarding it costs correctness. The research question of this paper is the review’s stronger form:

When can an agent stop remembering its execution history without losing the current evidence and obligations needed to finish correctly—and in which cost regime does the answer pay?

We answer it with a pre-registered, factorial mechanism study (Section 4). Seven retention policies—from full replay through observation masking, periodic summary/reset, bounded pages, and a verified board, to verified-frontier retirement with typed capsules, plus a demand-driven adaptive variant—are crossed with a scheduler factor in a full factorial, and consume byte-identical seeded work streams with temporal persistence, payload-bearing facts, persistent dependency footprints, and obligations that come due later. Accounting is split into construction, steady state, and a five-episode audit tail that compares each policy’s capsule against stream ground truth; two cost regimes are frozen ex ante, one byte-bound and one round-trip-bound. The v2 protocol and its five closure rules were committed to git before any v2 data was generated; the cohort and its tables were committed before any v2 manuscript text was written. Two protocol generations were run; the first’s design limitations are stated in Section 4.2 and its full record is retained in Appendix I.

The answer that survives both generations is a conversion rule with a regime-conditional payoff. *Necessity*: an agent may retire its history only after converting it into verified interface bindings, a typed pending-obligation set, and declared dependencies; policies that convert lossily—periodic summary/reset—fail the frozen acceptance contract at every scale measured (28/84 runs at the v2 scale; 21/63 at the v1 scale), silently reverting values to beliefs older than the working window and dropping obligations that cross reset boundaries. *Payoff*: the payoff is operations, not bytes. Frontier retirement saves only -8.0% on submitted bytes—the byte-level

result of the first generation, repeated—but cuts on-demand retrieval operations by 44.0% versus observation masking and reduces round-trip-regime cost per accepted run by 29.6% (gate $\geq 20\%$, sensitivity 17.8%–38.8% across the frozen round-trip constants), with the advantage persisting at four times the horizon (45.1% operations). The live-agent hypothesis this pre-registers is precise: savings, if any, come from avoided tool-call round trips, not from smaller prompts. *Audit*: the promoted mechanism passes a capsule audit against stream ground truth with zero defective service, and the audit’s development history includes a disclosed, test-caught reference fix (Section 4.4). *Rejected extensions*: a demand-driven adaptive-coordination variant was evaluated against a pre-registered promotion gate and rejected (-0.0% regime-O versus plain frontier; the gate required $\geq 10\%$), and the HACT certificate-tree evidence layout remains a conditional component—against a steel-manned flat ledger it wins dense updates and loses sparse ones, so it is selected by a measured-density rule with 0.0% regret rather than by default. The scheduler factor is cleanly separated: measured byte leak 0.000%, makespan reduction 87.5%.

Sections 5–11 develop the verification and evidence components the conversion rule relies on, unchanged in substance: immutable identities, packet cost models, layout learning, and the negative controls that bound what serialized evidence can claim. A capsule that carries verified bindings and pending obligations needs an evidence layout when its contents must cross a link, and the layout’s value is conditional in exactly the way the ablation measures. Sections 10–12 bound the delivery, adoption, and trust claims.

The negative controls remain: flat checker timings and per-episode verification overhead are limits of the serialization claim; the mechanism-level cohort is a limit of the retirement claim; the live-agent protocol of Section 12.1 has not been run, and no hosted model, token count, or billing figure appears anywhere in this paper as a measurement.

3 Related Work and Positioning

Context management for long-horizon agents. Observation masking with focused retrieval is the strongest simple baseline for bounded context [16]; context folding [17] branches and folds trajectories; periodic summary/reset is the generic compaction baseline. Our mechanism study steel-mans all three in a factorial design and reports where each fails: masking is near-optimal on submitted bytes but pays for every non-current read with a retrieval operation (44.0% more operations than frontier retirement at equal correctness), while summary/reset —cheapest on bytes—fails the acceptance contract on 56 of 84 runs through silent staleness and lost obligations. This locates our contribution precisely: not a better compaction heuristic, but a measured characterization of *which retained structures* (typed pending obligations, current verified bindings, dependency interfaces) make forgetting safe, and a pre-registered operations metric for the live cohort that could confirm it.

Multi-agent coordination. Decentralized shared context [19] and adaptive communication topologies address central-manager bottlenecks; controlled scaling studies [18] show coordination can hurt sequential or tool-heavy work. Our verified-board arm is exactly such a baseline and holds correctness (84/84) at the cost of shipping a full present-state row per gate; the scheduler factor is orthogonal and byte-neutral by design (measured leak 0.000%), consistent with the task-architecture alignment those studies report.

Testing, impact analysis, and incremental computation. Test prioritization seeks earlier fault exposure [1]; regression selection uses change dependencies [2, 3]. Our dependency footprints and completion records inherit that lineage; the certificate tree is an incremental-publication structure, not a selection structure, and its new ablation shows its value is conditional on update density (dense global updates favor the tree by 87.1%; sparse deltas favor a flat ledger).

Benchmarks. SWE-bench [14] and its harness [21] provide the external anchor the live cohort pre-registers in Section 12.1; no benchmark result appears in this paper because no live agent was run.

What is not claimed. Serialization savings are not compression savings: fewer padded bytes need not imply fewer compressed bytes, and we measure both where both exist. Flat ledgers and a trusted root-status endpoint can be much smaller than an archive of aggregation updates; they satisfy different delivery needs and are steel-manned, not hidden. No claim in this paper extends to agent speed, billing, or reasoning quality; those belong to the unrun live cohort, whose protocol is pre-registered with closure rules.

4 When Can an Agent Stop Remembering? A Pre-Registered Mechanism Study

4.1 The question

Long-horizon agents face a retention decision every episode: resubmit history, mask it and re-fetch on demand, or summarize it away. The research question is when the third option—*retirement*—is safe, and what it buys:

When can an agent stop remembering its execution history without losing the current evidence and obligations it needs to finish correctly, and in which cost regime does the answer pay?

The study is *mechanism-level*: every number below is produced by a deterministic engine that models information flow between a context policy and a seeded work stream. No hosted model is called. The answer that survives two frozen protocol generations is a **conversion rule with a regime-conditional payoff**:

An agent may retire its history only after converting it into (i) verified interface bindings for every fact another component will read, (ii) a typed pending-obligation set, and (iii) declared dependency interfaces. The payoff is not bytes: it is avoided retrieval round trips, which pay only in the round-trip-bound regime.

4.2 Study provenance and what changed between generations

Two protocol generations were run, both pre-registered.

Generation v1 (file-frozen 2026-09-23). Ten hand-picked arms over $3 \times 3 \times 7 = 63$ cells (630 runs) with three closure rules. Its closures fired: no byte-cost win for frontier retirement over masking (null closure), an attribution closure that forced the dimension-separated reading, and a conditional pass for grouped certificate layouts. v1’s protocol and engine were frozen in files before its cohort ran, but were committed to git only on 2026-09-24, after the v2 protocol commit; v1’s freeze is therefore file-attested only, and this ordering is disclosed here rather than smoothed over. The full v1 design, its numbers, and its deviations remain in Appendix I.

Generation v2 (git-frozen 2026-09-24, commit before data). v1 answered “what does each policy lose” but had four design weaknesses: a hand-picked arm list rather than a factorial, a single pooled byte threshold, capsule accounting evaluated only as an end-of-run total, and one cost regime. v2 repairs all four ex ante: (1) a full factorial over retention R (7 levels), scheduler

Table 1: Frozen v2 cohort at $S = \text{fixed}$: 84 cells $\times$ 7 retention policies. Regime-B bytes and regime-O cost are per accepted run; ops are on-demand retrieval operations per accepted run; audit defects (A3), staleness defects, and obligation losses are totals over the 84 runs of each policy. Acceptance is the frozen contract: any staleness defect or obligation loss rejects the run.

Policy	Accepted	Regime-B B/run	Regime-O B/run	Ops/run	A3	Stale	Lost
full-replay	84/84	12,007,957	12,007,957	0	0	0	0
masking	84/84	5,107,068	12,786,190	1,875	0	0	0
summary-reset	28/84	220,410	220,410	0	0	109,761	385
bounded-hier	84/84	789,534	5,748,522	1,211	0	0	0
verified-board	84/84	12,929,557	12,929,557	0	0	0	0
frontier	84/84	4,698,273	9,002,292	1,051	0	0	0
frontier-adaptive	84/84	4,698,350	9,005,733	1,052	0	0	0

S (2 levels), with factor contrasts replacing pairwise arm comparisons; (2) three-evaluation separation—every run’s accounting is split into construction (episodes 0–4), steady state (5–34), and a five-episode audit tail (35–39) in which an independent auditor compares the policy’s capsule against stream ground truth (A1 stale retention, A2 pending-set mismatch, A3 defective service); (3) two frozen cost regimes—bandwidth-bound B (bytes) and round-trip-bound O (bytes + 4,096 B per retrieval operation, with sensitivity at 1,024 and 16,384 B); (4) scale: four registry sizes ($n \in \{128, 512, 1,024, 4,096\}$), three families, seven seeds, forty episodes ($4 \times 3 \times 7 = 84$ cells $\times$ 14 policy runs = 1,176 runs), plus a long-horizon sub-cohort ($n = 1,024$, 160 episodes, 63 runs). Five closure rules were frozen before data generation; the v2 cohort and its tables were committed before any v2 manuscript text was written. One frozen-document miscount (“42 runs” for a grid the protocol itself defines as 63) was corrected ex ante in the engine notes and is disclosed here.

4.3 The frozen arms and closure rules

Retention levels. (1) **full-replay**: resubmit all history. (2) **masking**: identity-only context, every non-current read fetched on demand—the strongest simple baseline. (3) **summary-reset**: periodic summary/reset every 4 episodes with a window budget. (4) **bounded-hier**: bounded pages. (5) **verified-board**: a present-state row per known gate. (6) **frontier**: verified-frontier retirement—retain only the working set, boundary neighbors, dependency footprint, and pending obligations, as a typed capsule (160 B base, 16 B per retained binding actually held, 48 B per pending obligation, 64 B aggregate invariants). (7) **frontier-adaptive**: frontier plus a frozen demand-driven eviction rule (bindings unused for 8 episodes and opened ≥ 8 episodes ago are evicted outside the exempt base)—the *adaptive-coordination* candidate.

Closure rules (all frozen ex ante). *Null closure*: frontier is promotable only if it cuts regime-O cost per accepted run by $\geq 20\%$ versus masking pooled, with no family stratum regressing by more than 10%. *Attribution closure*: the scheduler factor is byte-neutral by design; a byte leak exceeding 1% of the grand mean voids all scheduler results. *HACT closure*: grouped layouts are promotable only as a conditional component—a $\geq 5\%$ win in the dense stratum plus a stated density decision rule with reported regret. *Adaptive-coordination closure*: **frontier-adaptive** is promotable only at equal acceptance, zero audit failure, $\geq 10\%$ regime-O reduction versus frontier, and no family regression above 5%. *Audit closure*: any arm with A3 defects in its audit tail is reported as audit-failing and cannot be promoted.

4.4 Results

The conversion rule is necessary: lossy conversion fails, typed conversion does not. Periodic summary/reset—the cheapest policy on submitted bytes—accepts only 28/84 runs:

Table 2: Family strata at $S = \text{fixed}$: regime-B bytes and regime-O cost per accepted run, and retrieval operations per accepted run, masking versus frontier. Global coupling leaves no operations gap (every read is a current fact); the operations advantage concentrates where reads cross the frontier.

Family	Masking B/run	Frontier B/run	Regime-O cut	Masking ops	Frontier ops
clustered	705,448	676,449	39.8%	230	77
independent	2,607,799	2,323,054	39.6%	5,394	3,075
global	12,007,957	11,095,317	7.6%	0	0

109,761 stale reads and 385 lost obligations at the new scale, the same structural failure v1 measured (30,820/280 on its smaller grid). The mechanism is exact: a summary reverts values to a belief older than the working window, and obligations crossing a reset boundary survive only through a lossy summary. Every policy that converts before retiring—frontier, and at higher cost the verified board and full replay—accepts 84/84 with zero staleness and zero loss. Bounded pages accept 84/84 but pay 1,211 ops/run in the steady phase; masking accepts 84/84 and pays 1,875.

The null closure passes in the round-trip regime, not the byte regime—exactly as pre-registered. Frontier retirement reduces regime-O cost per accepted run by 29.6% versus masking (gate $\geq 20\%$), with *no* family stratum regressing (worst stratum: none). In the bandwidth regime the same comparison is only -8.0%—we repeat v1’s honest negative and claim no byte-cost win. The two regimes separate because frontier’s advantage is *operations*: 44.0% fewer retrieval round trips per accepted run pooled (66.5% on clustered, 43.0% on independent; global coupling leaves no gap because every read is a current fact for both policies). Sensitivity across the frozen round-trip constant: 17.8% at 1,024 B, 29.6% at 4,096 B, 38.8% at 16,384 B—the advantage grows with the round-trip price and is positive across the entire pre-registered sensitivity range.

The audit tail is clean for the promoted mechanism. Across all 1,176 runs, the five-episode capsule audit records zero A3 defective service for every arm (audit closure: **clean**): retained bindings match stream generations, pending sets match the stream’s true pending sets in both directions, and no required read was served from defective state. The audit also caught a reference bug during development—the ground-truth pending set originally deduplicated obligation openings by first rather than latest opening, misclassifying correctly-held re-opened obligations as phantoms—which is disclosed in the repository history and fixed before the frozen cohort was regenerated.

Adaptive coordination was evaluated and rejected. The demand-driven eviction variant (**frontier-adaptive**) matches frontier’s acceptance (84/84) and audit cleanliness, but its regime-O reduction versus plain frontier is -0.0% (promotion gate $\geq 10\%$): the frozen use-clock rule evicts almost nothing the base frontier does not already evict, because the base exemptions (working set, obligations, footprint) already cover nearly everything the use clock would keep. Per the frozen closure it is **not promotable**, and the manuscript’s mechanism carries no adaptive component. This is the paper’s second honest negative after v1’s byte-regime null result, and it was decided by a pre-registered gate, not by narrative preference.

The scheduler factor is cleanly separated. The full factorial confirms byte-neutrality by design: the scheduler main effect on regime-B bytes is 0.000% of the grand mean (leak threshold 1%; attribution closure **not triggered**), while the makespan interaction is large—adaptive batching cuts makespan rounds by 87.5% (24,471 to 3,070 rounds/run pooled). Factor contrasts, not arm pairs, carry this claim: every $R \times S$ combination was run in every cell.

Table 3: Long-horizon sub-cohort ($n = 1,024$, 160 episodes, seven seeds, $S = \text{fixed}$). Ops and costs are per accepted run. The operations advantage persists at $4\times$ horizon; bytes stay comparable.

Family	Policy	Accepted	Ops/run	Regime-O B/run	Regime-B B/run
clustered	masking	7/7	745	4,449,648	1,398,713
clustered	frontier	7/7	284	2,559,028	1,394,594
clustered	frontier-adaptive	7/7	287	2,569,097	1,394,715
independent	masking	7/7	16,051	72,702,403	6,956,921
independent	frontier	7/7	8,943	42,826,402	6,194,704
independent	frontier-adaptive	7/7	8,958	42,888,791	6,196,238
global	masking	7/7	0	34,003,520	34,003,520
global	frontier	7/7	0	31,417,920	31,417,920
global	frontier-adaptive	7/7	0	31,417,920	31,417,920

Table 4: Evidence-layout ablation by family (pooled over sizes and seeds; held-out half). The density rule picks the tree exactly when it wins.

Family	Flat B/cell	HACT B/cell	Tree delta
clustered	150,386	204,631	36.1%
independent	277,129	344,823	24.4%
global	1,936,704	169,408	-91.3%

The operations advantage persists at four times the horizon. In the long-horizon sub-cohort ($n = 1,024$, 160 episodes), frontier cuts retrieval operations by 45.1% and regime-O cost by 30.9% versus masking at equal acceptance (63/63), while regime-B bytes remain within a few percent. The advantage does not decay with horizon; if anything the clustered stratum widens (66.5%→61.9% operations cut at the pooled gate, but 42.5% regime-O reduction versus 39.8% at the short horizon).

HACT remains a conditional component, now with a decision rule that costs nothing. The evidence-layout ablation (held-out half, order learned on the training half, flat steel-man at identical record cost) reproduces v1’s shape at the new scale: the learned tree wins the dense global stratum (169,408 vs 1,936,704 held-out B/cell, an 87.1% cut) and loses the sparse strata. The frozen density rule—select the tree iff measured training-half update density ≥ 0.5 —matches the per-cell oracle on every one of the 84 cells (0.0% regret). HACT therefore stays what v1’s closure made it: a conditional serialization component with an observable selection signal, not the paper’s center. Sections 6–11 retain the component development with its scope labels unchanged.

4.5 What the mechanism study does and does not establish

The conversion rule and the regime-conditional payoff are structural predictions about policies, not measurements of agents. The audit checks a policy’s capsule against stream ground truth in a deterministic engine; a live agent’s summaries, retrievals, and obligations are not instrumented by it. The 4,096 B round-trip constant is a frozen accounting choice, not a measured latency; the sensitivity range brackets it but does not replace a live measurement. Byte numbers are the frozen accounting of a frozen model of information flow; no hosted model was called, no token or billing claim is made, and the BPE and physical-WAN gaps of Section 11 remain open. The live cohort of Section 12.1 pre-registers operations as its primary cost measure precisely because that is where the mechanism result predicts the payoff.

5 Model and Trust Boundary

The retirement study of Section 4 asks when history can be discarded; this section and the two that follow define the verification substrate that makes the discarded history reconstructable and auditable—the stable identities, generations, and evidence layouts a typed capsule refers to. These component results carry their v5.1 scope labels unchanged; they are no longer the paper’s center, and the new ablation (Table 4) now bounds when their layout matters.

5.1 Semantic registry, execution, and display

Let $\mathcal{R} = (g_0, \dots, g_{n-1})$ be a nonempty approved registry of unique check IDs. The candidate identity is

$$e = (\text{epoch}, \text{snapshot}, \text{checker}, \text{environment}, \text{registry}). \quad (1)$$

The snapshot covers source and protected acceptance inputs. The checker and environment are explicitly bound; an epoch increases even after an ABA return to identical source bytes. Only a trusted runner submits final PASS, FAIL, or UNKNOWN records. Missing, stale, skipped, incompatible, or timed-out outcomes are not passes.

Immutable identities versus mutable execution state. The model splits its state into two disjoint layers. *Immutable check identities* are the registry $\mathcal{R}$, the candidate identity e , and the acceptance predicate each check declares; they are fixed at publication and are never rewritten by execution, layout choice, or learning. *Mutable execution state* is everything a run may change: the result tokens, their generation counters, epoch-local timestamps, and the execution permutation σ_e that determines when checks run. A record’s identity is its semantic ID bound to e ; its display position is not part of that identity. This separation is what makes evidence serialization under fixed schedules well-posed: the schedule and outcomes are observed inputs, and only the aggregation layout is chosen. No layout or learner writes execution state, and no execution mutates the registry or acceptance predicates.

An execution permutation σ_e determines when checks run. A layout $L = (\pi, T)$ independently maps tree positions to semantic registry indices. For a node covering positional interval $[i, j]$, its semantic coverage is

$$G_L(i, j) = \{g_{\pi(i)}, \dots, g_{\pi(j)}\}. \quad (2)$$

The stored layout must contain an exact permutation, and every internal node’s children must partition its positional interval without gaps or overlaps. The task graph is outside this mathematical tree; one worker can discharge many checks, and one integration check can cover several workers.

The acceptance predicate is independent of both permutations:

$$A(e) = \mathbf{1}\{\text{every } g_i \in \mathcal{R} \text{ has trusted, matching, current PASS evidence for } e\}. \quad (3)$$

An observed fail permits rejection before all checks finish. It does not assert coverage of the unexecuted remainder. Acceptance is relative to the approved predicates, not proof of all program behavior.

5.2 Packet and lifecycle costs

Each internal node of arity d emits a fixed-width packet

$$c(d) = H + Sd, \quad 2 \leq d \leq b = \lfloor (C - H)/S \rfloor. \quad (4)$$

Here $H = 512$, $S = 320$, and $C = 3072$ bytes, so $b = 8$. A mandatory field that exceeds its slot is rejected, not silently truncated. The cap applies to one certificate packet, not to all

project memory, an entire registry manifest, or every LLM prompt. A manifest is a referenced control-plane artifact; its measured size is charged on migration.

Let $W_{e,1}, \dots, W_{e,q_e}$ be the semantic IDs whose final records become available between successive refreshes. For a strict new candidate epoch, the observed cost is

$$\ell_e(L) = W(L) + \sum_{a=1}^{q_e} \sum_{v \in \text{int}(T)} c(d_v) \mathbf{1}\{W_{e,a} \cap G_L(v) \neq \emptyset\}, \quad W(L) = \sum_v c(d_v). \quad (5)$$

The first term constructs the full initially unknown tree. The CLI batches an executed prefix into groups of eight, flushing its final partial group. A successful strict epoch runs the whole registry. This conservative implementation does not reuse a pass across distinct source snapshots.

We separately study persistent-state refresh traces that omit $W(L)$ at each tick, as in a cache with correctly specified incremental dependencies. Each topology installation is still cold and explicitly priced. These traces evaluate the controller’s cost model; they are not the strict candidate-gate timing experiment. They do not demonstrate safe dependency inference from arbitrary code.

5.3 Why more than invalidation or pairwise similarity is needed

A classical result shows that marginal check activation cannot identify the best aggregation tree under joint invalidation (Appendix E). A further result shows that even the full initial invalidation distribution cannot identify optimal progressive verification: early rejection and completion batching change the update sequence. Both distinctions remain relevant after allowing permutations.

Proposition 1 (Pairwise affinity is insufficient for full-wave scoring). *There exist two distributions with identical singleton and pairwise co-activity probabilities, hence identical Jaccard affinities, but different activation probabilities for the same three-check subtree.*

Proof. On IDs $\{0, 1, 2\}$, let P be uniform on $\emptyset, \{0, 1\}, \{0, 2\}, \{1, 2\}$ and Q uniform on $\{0\}, \{1\}, \{2\}, \{0, 1, 2\}$. Every singleton has probability $1/2$, every pair has joint probability $1/4$, and its Jaccard affinity is $1/3$. Yet the probability that the triple is hit is $3/4$ under P and 1 under Q . Thus a packet emitted on that interval has different expected cost. $\square \quad \square$

This elementary parity construction is not a new fact about higher-order statistics. Its implication for the compiler is specific: a pairwise clustering objective cannot substitute for scoring recorded completion hyperedges. It does not, by itself, prove that these two distributions have different globally optimal layouts.

6 Learning Layouts Without Changing the Contract

6.1 Practical catalogue and full-wave selection

Training waves define counts N_i , N_{ij} and affinity

$$a_{ij} = \frac{N_{ij}}{N_i + N_j - N_{ij}}, \quad (6)$$

with zero for an unobserved union and zero diagonal. The practical proposal uses established average-linkage clustering: join the closest current clusters, choose endpoint orientation by cross-boundary affinity, and break ties canonically. All semantic IDs, including disconnected and unseen checks, remain included exactly once.

Given the learned order, group consecutive leaves into blocks of at most eight, promote a singleton unchanged, and repeat upward. This bottom-up compact balanced construction does not call a tree optimizer. The catalogue retains the incumbent first, includes original-order fixed8 when distinct, and proposes learned-order fixed8. Identical layouts are deduplicated. A separate validation cohort scores complete completion waves, including initial construction; ties retain the incumbent. The API returns an advisory choice, not deployment authorization. The registry and packet cap cannot change during selection.

For any proposal π , training episodes supply interval weights

$$r_\pi(i, j) = 1 + \frac{1}{m} \sum_{e=1}^m \sum_a \mathbf{1}\{W_{e,a} \cap G_L(i, j) \neq \emptyset\}. \quad (7)$$

These weights also define the exact-DP reference below. Pairwise affinity proposes an order; it does not replace the full-wave cost. The practical catalogue contains at most three candidates and uses no DP. Tree construction is $O(n)$ after ordering, not a linear-time claim for affinity construction, clustering or validation scoring. The reference implementation retains spectral ordering and exact fitting for research comparisons; neither is invoked by the practical default.

6.2 Exact conditional tree synthesis

Let $D(i, j)$ be the optimal subtree cost and $F_k(i, j)$ the optimal cost of an ordered forest with k nonempty children. With infeasible states equal to $+\infty$,

$$D(i, i) = 0, \quad (8)$$

$$F_1(i, j) = D(i, j), \quad (9)$$

$$F_k(i, j) = \min_{i+k-2 \leq u < j} \{F_{k-1}(i, u) + D(u+1, j)\}, \quad (10)$$

$$D(i, j) = \min_{2 \leq k \leq \min(b, j-i+1)} \{r_\pi(i, j)c(k) + F_k(i, j)\}. \quad (11)$$

Theorem 2 (Exactness conditional on order). *For a fixed π , nonnegative interval weights, and the packet model in Equation (4), Equation (11) returns a minimizing feasible tree in $O(bn^3)$ time and $O(bn^2)$ space.*

Proof. Every feasible root has an arity and a contiguous partition of its interval. Its local contribution depends only on that interval and arity; all remaining contributions belong to independent child subtrees. The forest recurrence enumerates every such partition. Induction on interval length proves exactness. There are $O(bn^2)$ states, each scanning at most n splits. The apparent $F_1 = D$ dependency is used only on shorter intervals when constructing a forest with at least two components. $\square$ $\square$

This is a specialization of classical alphabetic-tree optimization [6]. The historical three-proposal reference multiplies compilation time by a constant, not by $n!$. Dense spectral computation can itself cost $O(n^3)$; the reference clustering implementation uses $O(n^2 \log n)$ heap operations and $O(n^2)$ storage.

Proposition 3 (Validation after learned ordering). *Condition on a training set that fixes K complete layouts. For m independent identically distributed validation epochs with each layout cost in $[0, B]$, let $\hat{L}$ minimize empirical validation cost. With probability at least $1 - \delta$,*

$$J(\hat{L}) \leq \min_{L \in \mathcal{C}} J(L) + 2B \sqrt{\frac{\log(2K/\delta)}{2m}}. \quad (12)$$

Algorithm 1 Deployable evidence control; inference and checking remain separate**Require:** Protected registry, checker and snapshot identities; frozen layout catalogue

```

1: for each event do
2:   Choose from past evidence using Equation (14), or keep the frozen layout
3:   if layout changes then
4:     Revoke tickets, increment generation, remap stable IDs, rebuild and price the installation
5:   end if
6:   Receive trusted completion records for the fixed check schedule
7:   Reject stale, conflicting or malformed records; aggregate bounded packets
8:   Publish only the complete current contract result; otherwise retain FAIL/UNKNOWN
9:   Expose a root-status card; retrieve diagnostic packets only when needed
10:  Score the observed event and append its cost vector to the bounded window
11: end for

```

Proof. Conditionally, the catalogue is fixed. Hoeffding’s inequality and a union bound over K layouts bound every empirical mean error by the displayed radius without its factor two. Inserting empirical costs on either side of the empirical minimizer yields the result. $\square$ $\square$

The guarantee is only relative to the fixed catalogue. It neither justifies choosing among learned permutations on their training errors alone nor covers distribution shift. We do not report this potentially loose bound as an empirical confidence interval; reported intervals use independent experimental seeds.

6.3 Compiler identity and scope

The practical default has no hidden microtree optimization. Exact DP is a reference and optional offline candidate. The block-backbone proposition and implementation remain in Appendix D solely as an ablation; its restricted-class exactness does not override the observed simpler-baseline advantage. At 1,024 checks, order learning still takes seconds in our measurement, so layouts should be cached and compiled away from the agent’s per-action path. Validation of padded costs does not optimize compressed archives or model tokenization; adoption must measure the actual delivery representation.

7 Practical Online Selection: Price the Window, Not the Promise

Freeze a compiled catalogue $\mathcal{C} = \{L_1, \dots, L_K\}$. At event t , choose a layout before observing the current completion wave; afterwards score that same wave under every catalogue member. This is full-information feedback only because the check schedule is held fixed. It is not counterfactual feedback for arbitrary agent behavior.

A different layout is installed by serializing its canonical manifest and reconstructing its complete certificate tree:

$$M_{ij} = \begin{cases} 0, & L_i = L_j, \\ |\text{manifest}(L_j)| + W(L_j), & L_i \neq L_j. \end{cases} \quad L_T = \sum_t \ell_t(A_t) + \sum_{t=2}^T M_{A_{t-1}, A_t}. \quad (13)$$

The initial installation is separate. These byte prices match the evaluated uncompressed migration protocol. A deployment using compression must measure its own priced wire representation; the DP objective is not automatically an optimum for a compressed log.

The primary practical rule. Let $\bar{\ell}_{t,i}$ be the mean cost on the last $\min(w, t - 1)$ observed events, and $j_t = \arg \min_i \bar{\ell}_{t,i}$ with deterministic ties. Starting from the calibrated incumbent a , use the frozen rule

$$A_t = \begin{cases} j_t, & w(\bar{\ell}_{t,a} - \bar{\ell}_{t,j_t}) > M_{a,j_t}, \\ a, & \text{otherwise,} \end{cases} \quad w = 64. \quad (14)$$

With no history, keep a . A tie in estimated savings does not trigger migration. During warm-up, the forecast horizon remains w even if fewer events were observed, matching the original baseline. This is an empirical break-even forecast, not proof that past savings repay a future switch. It can lose under abrupt reversal; preserving the frozen layout is always a legitimate alternative.

The runtime API enforces `choose()` before `observe()`, bounds history to 64 vectors, rejects nonfinite costs, and checkpoints only at event boundaries. Replaying all 80 archived streams, with periodic checkpoint/resume, must reproduce the original baseline’s full action and byte-cost paths exactly. The exercise validates implementation equivalence, not a new performance dataset.

Why fixed-share is no longer primary. On the archived stationary, long-shift and rapid-shift streams, the coupled fixed-share policy is more expensive than the priced window. Its standard tracking analysis therefore moves to Appendix C, with all negative results retained. We do not manufacture an adversarial workload merely to make that policy win. Neither the appendix bound nor the empirical window result establishes general superiority over other switching-cost policies.

8 Stable Evidence Through Permutation and Migration

A result token carries (epoch, g_i, b_i) , where b_i binds the semantic ID to the registry, source snapshot, checker, and environment. It does not carry a display position as its identity. A layout-specific kernel maps stable IDs to current positions. During migration it rebuilds from canonical trusted records; it never transplants an ordinal cache entry into a different permutation.

Publication tickets additionally carry a monotone layout generation and layout hash. A layout switch revokes the issued ticket, advances the generation, and forces reconstruction. Returning to an earlier tree therefore cannot revive that tree’s old ticket. An in-flight token may remain valid through a layout-only change, but a new candidate epoch fences it. Duplicate identical evidence is idempotent. Conflicting final records in one epoch revoke authorization: a recorded fail remains a fail; otherwise the obligation becomes unknown. Resolution requires a new epoch, not a more persuasive agent summary.

Theorem 4 (Layout-independent local publication soundness). *Assume complete approved predicates and dependency identities, a trusted checker executing immutable identified inputs, collision-resistant binding, and atomic local epoch/generation checks. For any sequence of valid layout permutations, any advisory learning outputs, and any delivered evidence order, authorization implies Equation (3). If each checker is sound for its declared predicate, those predicates hold on the authorized candidate.*

Proof. A canonical pass requires a current token for the same semantic ID and immutable identities. Permutation is a bijection, so rebuilding moves each such record to exactly one position without changing its meaning. Exact child partitioning preserves coverage and pass/fail/unknown counts by induction up the tree. The root can be authorized only with $(n, 0, 0)$ and matching current registry, epoch, generation, and layout. Old tickets fail those checks. Conflicts prevent authorization. No choice made by the learner writes trusted outcomes or changes the registry. $\square$ $\square$

This is a conditional local guarantee, not remote attestation or distributed consensus. An incomplete suite, compromised checker, omitted dependency, or nondeterministic external service

Table 5: Separate ledgers avoid converting byte savings into unsupported latency or token claims.

Measurement	How obtained	In optimized objective?
Certificate bytes	Actual padded packets; independently reconstructed from waves	Yes
Migration bytes	Canonical layout manifest plus forced full rebuild	Yes, online traces
Fit and scoring time	Monotonic-clock measurements	Reported separately
Preparation, copying, checking, aggregation	Wrapper timers and outer call duration	Reported separately
Active harness episode	Supervisor startup through completion and shutdown	Reported separately
Prompt bytes and construction time	Captured actual HTTP messages and context builder	Reported; not tokens
Hosted token usage and API billing	No provider inference was run	Unknown, never zero-imputed

can invalidate a broader correctness claim. The CLI’s file copy is not an OS sandbox. An external merge/deployment service must atomically compare the verified candidate and integration-head identities at commit; a report file alone is not that transaction.

9 Implementation and Cost Accounting

The artifact adds `practical.py`, `deployment.py`, `token_budget.py` and `tls_probe.py`. The practical CLI produces a balanced catalogue without calling either exact or block DP. The historical exact and block compilers remain separately callable. The existing stable-ID guard, priced window, codec and Future Code bridge are retained. The strict CLI compares protected input sets and hashes, detecting added tests or configuration files as well as changed and deleted ones. Schema-2 contracts bind the approved checker and environment and refuse silent reuse of legacy contracts. Source manifests are checked before and after snapshotting and verification; mutation of an observed snapshot prevents a pass.

The bridge runs as an ordinary subprocess acceptance gate in the unmodified Future Code v1.2 runtime. The supervisor still owns worker leases, private attempts, bounded contexts, and serialized integration. Non-pass exits cause candidate rejection. Detailed evidence and binary certificate packets live outside candidate write scopes; only a bounded JSON result card returns to the worker. The plug-in accepts a precompiled layout. The online selector and migration guard are separately integrated and stress-tested as a library path; a live, self-retuning hosted deployment was not evaluated.

A different scalar objective could combine these quantities only with declared unit conversions and observable counterfactual costs. This implementation does not optimize unknown model latency or billing. The exact-byte claim also depends on the stated padded format; compression defines a different, potentially nonseparable objective for which this DP is not claimed exact.

10 Service-Matched Delivery and Admission

10.1 Do not conflate three consumers

The local aggregator receives bounded packets; a logging collector may receive their complete stream; a supervisor normally receives a status projection and retrieves details selectively. Define $\psi(e, c)$ as canonical status text containing semantic identity, counts and local authorization, but

no display positions. For the same trusted semantic report,

$$\psi_{L_1}(e, c) = \psi_{L_2}(e, c). \quad (15)$$

Its byte length and token count under any fixed deterministic tokenizer are therefore identical. This identity is not a new theorem about tokenization. It supplies a negative control: layout savings cannot reduce a layout-independent root input. The card is a projection of a trusted report, not authority to accept arbitrary worker JSON.

Diagnostics preserve full header/card records and page whole packets under a 16,384-byte limit. A packet that cannot fit is rejected, not truncated. For a real supplied tokenizer τ and explicitly reserved framing budget r , the new admission API additionally requires

$$|p|_{\text{UTF8}} \leq C_{\text{bytes}}, \quad |\tau(\text{decodeUTF8}(p))| + r \leq C_{\text{tokens}}. \quad (16)$$

The implementation retokenizes each full proposed page, including its prefix. Token counts are not assumed additive across concatenations. Greedy packing guarantees the tested limits, not a minimum number of pages. Toy counters in API tests are not BPE measurements. The optional measurement command requires actual tiktoken software and vocabulary [20], fingerprints the encoding, and records per-page content hashes. Missing encodings produce UNKNOWN and null counts; there is no byte-ratio fallback. Provider framing and billed usage require a separately measured adapter, not a guessed constant.

10.2 Encoding and the simplest sufficient service

Canonical unpadded JSON, independently compressed packets, and whole-epoch DEFLATE preserve their input records with different operational boundaries. Whole-epoch compression requires the completed archive, whereas per-packet compression permits incremental export. The padded objective is not an exact optimizer of either compressed bytes or tokens. Practical selection is advisory until its gain is measured with the deployed codec.

For a trusted consumer needing only final status, a root endpoint or flat ledger can be much cheaper than the intermediate stream. This changes the diagnostic service, not necessarily the underlying check set. Comparing that smaller service with an exhaustive stream cannot demonstrate that a tree optimizer is efficient or necessary. The strict checker, data representation, and delivery demand must each be named explicitly.

10.3 Adoption cost is not marginal layout cost

Consider q copies of the same workload and delivery service through one nonoverlapped bottleneck. Let B_0, B_1 be codec-matched application bytes per copy, $R > 0$ effective application goodput, $E \geq 0$ additional encoding/decoding work per copy, and $C \geq 0$ consistently amortized one-time incremental work. Common setup cancels only when actually shared. The conditional net benefit is

$$d = (B_0 - B_1)/R - E, \quad \Delta(q) = qd - C. \quad (17)$$

A strictly profitable integer copy count exists only when $d > 0$, and then

$$q_{\min} = \lfloor C/d \rfloor + 1. \quad (18)$$

This is accounting, not a new networking or regret theorem. A zero gain retains the baseline. The implementation uses decimal arithmetic at equality and rejects mismatched service IDs. Missing C , E or R returns UNKNOWN, not zero. An explicit zero is permissible only as a declared assumption or measurement.

When both variants already use the same protected gate, C is extra layout work, not the full strict-wrapper cost. When adopting strict verification from ordinary testing, its additional

cost must be justified independently; the services differ and the planner returns `INCOMPARABLE`. Repeating exports amortizes one candidate’s preparation, not the preparation of future new candidates. Parallel links, computation overlap, cached diagnostics, queueing and different connection policies require direct critical-path measurement. High RTT alone does not make byte savings larger if both exports pay the same setup and request count.

10.4 Remote measurement and security boundaries

The historical transfer experiment is receiver-paced loopback, not a WAN. The new opt-in measurement kit uses verified TLS server certificates and hostname checking; upload and nonlocal bind require explicit switches. Frames and decompressed archives are bounded, and acknowledgments check received length and digest. Local tests exercise cold and reused TLS sessions, rejected hostnames, untrusted roots, malformed acknowledgments and truncated records.

These tests validate a protocol path only. They provide no physical-WAN throughput, jitter, packet-loss or retransmission measurement. The supplied protocol records host and path provenance, codec, connection policy and matched payloads, and requests separately observed RTT/goodput and transport diagnostics where available, consistent with the distinctions in TCP throughput testing [22]. Fields that the minimal client cannot observe remain null. There is no receiver-side artificial pacing in a qualifying WAN trial. No public upload, account creation or credential search was performed.

TLS authenticates the configured server, not the checker’s truthfulness. Protected reference data, immutable candidate identity and local monotone state remain required for Theorem 4. There is no remote attestation, Byzantine consensus, client authorization service, non-repudiation or OS sandbox. Replacing both evidence and its trusted reference still defeats a digest-only comparison. The receiver is a bounded measurement utility to use behind an operator-controlled firewall, not a production ingestion service.

11 Evaluation of the Practical Revision

11.1 Provenance and frozen questions

The v5 contract was saved before the new runs. It fixes three sizes (128, 512, 1,024), three completion families (clustered, independent, global), and three seeds (55011–55013). Each of the 27 workloads has 128 training epochs, 128 validation epochs and 256 held-out epochs. Random draws are sequential disjoint samples; deterministic global observations naturally coincide. The catalogue is fit only on training data, selected on validation, and reported on held-out data. Cold strict-epoch construction is included in every objective. The nine size/family summaries use seed means and observed ranges, not episode-pooled significance tests.

A separate study re-aggregates twelve archived source candidates under three layouts, producing 36 layout replays. All use the same original-order trusted check records, the same executed sequence and batches of eight. Average-linkage orders were learned on the earlier training cohort and frozen before these candidates. None of these replays is a fresh checker execution, independent code repair, or new agent trial. Historical v3/v4 data remain in the appendices with their original cohort identities.

11.2 Practical default without exact or block DP

On clustered completions, mean held-out reductions are 23.43%, 10.75% and 5.56% as registry size increases (Table 6). Global updates give no reduction; independent updates give zero except for one of the three 512-check seeds, whose small gain produces a 0.07% mean. These controls do

Table 6: New v5 synthetic study. Reduction is selected-layout held-out padded cost relative to the original-order balanced incumbent; positive is better. Fit includes affinity/order learning and catalogue construction, not validation scoring. Three seeds per row; ranges are descriptive.

Checks	Family	Mean (%)	Range (%)	Fit median (s)	Changed
128	cluster	23.43	22.40–24.45	0.018	3/3
128	independent	0.00	0.00–0.00	0.016	0/3
128	global	0.00	0.00–0.00	0.032	0/3
512	cluster	10.75	10.06–11.13	0.523	3/3
512	independent	0.07	0.00–0.20	0.548	1/3
512	global	0.00	0.00–0.00	0.645	0/3
1,024	cluster	5.56	5.16–5.80	2.991	3/3
1,024	independent	0.00	0.00–0.00	3.086	1/3
1,024	global	0.00	0.00–0.00	3.462	0/3

Table 7: New serialization of twelve archived source candidates, four per project. The first three columns compare learned-order balanced8 with original-order fixed8. The last column compares learned balanced with learned exact after epoch compression: positive is a penalty. All include the codec’s application framing.

Project	Padded (%)	Canonical (%)	Epoch (%)	vs. exact penalty (%)
networkx	48.63	48.74	41.84	11.68
toolz	8.78	8.70	4.36	1.52
future	29.12	29.23	18.53	6.67

not support universal value from ordering. The selected method includes no DP, and a regression test makes any call to the exact fitting helper fail on this path.

The full catalogue fit median at 1,024 checks remains about 3 seconds; order learning is not free. The fixed construction can reuse a previously approved order in linear work, but no linear-time claim is made for the full learned pipeline. The smaller percentage at larger registries is consistent with substantial layout-independent cold construction in this benchmark; it should not be interpreted as a learned scaling law.

An initial execution exceeded its 120-second tool limit after saving 26 complete records. No process remained; explicit resume skipped those exact records and completed the remaining frozen seed without changing data or selection. All 27 observations and the interruption record are retained.

11.3 Source-evidence replays: simplicity is not byte optimality

Learned balanced layouts reduce whole-epoch compressed bytes by 41.84% for NetworkX, 4.36% for Toolz and 18.53% for Future Code (Table 7). These are not the larger v4 exact-layout savings relabeled as results of the new default. The balanced representation remains larger than learned exact in every project here. Exact fitting may therefore be appropriate when amortized delivery justifies its additional compilation; the practical default trades that potential gain for simpler fitting.

Across the twelve cases, batch-compressed application bytes are 416,046 for fixed8, 299,081 for learned balanced and 278,831 for the re-serialized learned exact layout. Semantic counts, verdicts and local authorization agree across all 36 replays. Their provenance binds the archived check records by digest. Serialization includes newly bound layout/run metadata; numerical bytes need not exactly reproduce a prior archive’s compressibility. The independent audit checks actual emitted bytes rather than substituting earlier totals.

Root-status text is identical in all twelve triples and occupies 593–594 UTF-8 bytes. Exhaustive

Table 8: Conditional minimum profitable copies of the same twelve-candidate cohort with whole-epoch compression. $E = 0$ is an assumption. The columns assume extra cost per candidate, paid once; cohort costs are 0.6, 6 and 60 seconds. These are not measured strict-wrapper overheads.

Assumed goodput	0.05 s/candidate	0.5 s/candidate	5 s/candidate
8 KiB/s	1	1	5
64 KiB/s	1	4	34
1 MiB/s	6	54	538

diagnostic pages total 148 for fixed8, 95 for learned balanced and 87 for learned exact. The new default therefore uses 35.81% fewer pages than fixed8, not v4’s 41.22% exact-layout reduction. The largest emitted page is 16,364 bytes under a 16,384-byte limit. Reading every page is a specified exhaustive diagnostic service, not observed LLM behavior. There is no measured token reduction; root token equality follows from identical text for a fixed tokenizer, but no numerical BPE counts are asserted.

11.4 Codec-matched deployment scenarios, not measured speedups

The measured fixed-to-balanced compressed byte difference is 116,965 per twelve-candidate cohort. Equation (17) gives 1.784744 seconds of gross delivery saving per cohort copy at an assumed 64 KiB/s goodput, with $E = 0$. This is arithmetic from replayed byte sizes, not the older 2.21-second measured loopback saving of a different layout.

The 5-second column (Table 8) is a deliberately declared sensitivity scenario, not an estimate transferred from the historical harness. At 64 KiB/s and one export per candidate, the assumed 60 seconds of cohort overhead overwhelms the approximately 1.785-second delivery saving. Repeated delivery might change the equation, but the application must actually require it. With new candidates on each repetition, preparation recurs and cannot be amortized as the same C . All 27 scenarios (three codecs, three rates, three overhead assumptions) are available, including negative outcomes. The planner rejects mismatched services and missing measurements rather than inventing a benefit.

11.5 Historical controls remain controls

The earlier 288 loopback transfers belong to v4, use application pacing rather than real WAN conditions, and are not rerun or relabeled as v5 trials. For twelve compressed fixed/exact exports at 64 KiB/s, their measured totals were 6.624 and 4.413 seconds. The larger uncompressed totals were 33.499 and 18.954 seconds. Compression of fixed8 was more impactful than selecting an uncompressed learned layout. None of those durations is an IP/TLS throughput measurement or predicts a loss-prone path.

The original 39.69/39.77-second checker totals and 12.98/18.11/18.12-second median direct/fixed/learned harness durations remain negative controls. They came from different source/harness fixtures, so this paper does not subtract their times across cohorts. The historical window-controller and block-compiler evidence is retained unchanged. Block DP’s 1.54–5.34% penalty to same-order balanced8 is the reason it is now an ablation, not a contribution advertised as deployment acceleration.

11.6 Unmeasured endpoints and tested interfaces

Both requested BPE encodings, `c1100k_base` and `o200k_base`, remain UNKNOWN: the software was absent, package/vocabulary fetches failed, and a connector-retrieved build artifact was expired. No substitute tokenizer, model alias or byte division was used. The measurement tool records

exact content-only counts when a real encoding is available; its unit tests use explicitly labeled toy counters only to validate admission logic.

No two authorized remote measurement hosts or designated live model endpoint were available. Physical-WAN and hosted-LLM trial counts are zero. The new TLS utility’s local compatibility tests check transport integrity, certificate rejection and bounded framing; they do not measure WAN jitter/loss or close the agent-evaluation gap. A reproducible deployment protocol is supplied, rather than a claimed remote experiment.

11.7 Software and independent audit

The v5 prototype passes 329 tests, including 76 new cases; the existing version-equality assertion is updated to require 5.0.0 consistently in runtime and package metadata. A duplicate test-module basename and the initially inconsistent version were corrected without weakening behavior checks; failure logs are retained. New tests cover no-DP compilation, service mismatch, missing costs, exact break-even, whole-text token admission, TLS trust failures and deliberate data tampering.

A separate auditor imports no producer optimizer, selector or codec. It reconstructs full-initialization and wave costs for all 27 workloads, checks 36 serialized replays against original evidence digests and verdicts, reconstructs all codec byte sizes and diagnostic pages, and independently checks 27 planning scenarios and 131 summary comparisons. These checks validate recorded arithmetic and provenance, not third-party security certification. Release verification separately records extracted-source tests, the unchanged Future Code suite, installed-source hashes and archive integrity.

12 Deployment Recommendation, Limitations, and the Pre-Registered Live Cohort

12.1 Deployment recommendation

The recommendation now has two tiers, one measured at mechanism level and one still protocol-only.

Mechanism-level (measured this paper). Where a service already requires protected verification evidence and a constrained channel, the serialized-evidence component of Sections 5–11 applies as before. The retirement study adds a sharper rule for context policy: a policy may retire history only after converting it into current verified bindings, a typed pending-obligation set, and dependency interfaces; policies that rely on lossy summaries of obligations, or that revert values to beliefs older than the current working window, measurably fail at every scale tested (Section 4.4). The conversion is necessary; whether it is worth adopting depends on the cost regime: it buys 44.0% fewer retrieval round trips at equal correctness, which is worth a great deal when round trips are expensive (regime-O cost 29.6% lower than masking) and almost nothing when only bandwidth is priced (regime-B bytes -8.0%). An adopter should therefore price their retrieval path before choosing: the frozen 4,096 B-per-round-trip accounting brackets the answer with a sensitivity range (17.8%–38.8%), not a measurement. Between grouped certificate layouts and flat delta ledgers, select by update density: the measured-density rule matched the per-cell oracle on every cell of the frozen cohort (0.0% regret); dense updates favor the tree, sparse deltas favor the flat ledger, and the density signal is observable on a training half of any real stream. Do not adopt the demand-driven adaptive variant: its pre-registered promotion gate failed (−0.0% versus plain frontier).

Live-agent cohort (pre-registered, not run). The mechanism study predicts, but cannot confirm, that avoided retrieval operations translate into live savings. The confirmatory cohort is therefore pre-registered with operations as the primary cost measure:

1. **Workload.** Multi-stage changes on a fixed public issue benchmark (SWE-bench-class), screened for dependency coupling; the same frozen independent acceptance contract for all arms; equal total budgets, sandbox, and protected state; every failed branch in the numerator.
2. **Arms.** The arms the mechanism study separates: observation masking; summary/reset; verified-frontier retirement; and the density-selected evidence layout (flat delta ledger versus grouped tree, chosen per the measured update density of the pilot). The demand-driven adaptive variant is excluded because its promotion gate failed at mechanism level; including it would re-litigate a pre-registered decision.
3. **Primary outcomes.** Tool-call round trips per accepted project (the pre-registered live-cost measure, from the mechanism result); cost per accepted project (all provider metering preserved, unknowns never zero-imputed); acceptance rate; deadline success with censored timeouts. Pre-specified noninferiority margin on acceptance: no arm may trail masking by more than 5% absolute.
4. **Closure rules.** *Null closure:* if frontier retirement does not reduce tool-call operations per accepted project by at least 20% versus masking at noninferior acceptance, the live-cohort hypothesis is rejected and retirement is not promoted. *HACT closure:* if the density-selected layout does not beat flat-delta evidence by at least 5% on measured evidence bytes at equal correctness, the tree is not deployed.

Two explicit closure rules prevent indefinite iteration; either rule terminates the program with a written result. This protocol is frozen here, before any live data collection.

What remains unmeasured. No hosted model was called for any number in this paper; BPE tokenization, physical-WAN behavior, and provider billing remain unknown and are never zero-imputed. The mechanism cohort’s cost model is a frozen accounting of a frozen model of information flow; real payloads are not lognormal, and real summaries may retain obligations better or worse than the seeded boundary survival. The round-trip regime’s 4,096 B constant is an accounting choice bracketed by a sensitivity range, not a measured latency. The tool-call prediction is the only bridge from the mechanism result to live cost, and it is pre-registered as such rather than assumed.

Trust and scope. Future Code remains a single-host supervisor. The guard and TLS utility are not a distributed consensus protocol or fault-tolerant deployment. A passing declared suite cannot establish absence of arbitrary defects. Authorization requires an uncompromised checker, correct protected reference data, immutable candidate identity, and local monotone state; Theorem 4 stands as before. TLS authenticates the configured server, not the checker’s truthfulness. The receiver is a bounded measurement utility behind an operator-controlled firewall, not a production ingestion service.

A Audit and Reproduction

A.1 Exact data generation

The synthetic generator uses NumPy’s `default_rng` initialized with a `SeedSequence` whose entries are `[260918, n, seed, stream]`. Stream zero produces training episodes and stream one held-out episodes. Group construction uses `SeedSequence([7001, seed, n])`. Pilot seeds 1000

and 1001 are excluded from all main tables. Scaling sizes other than 128 have their own saved records; the four reused 128-gate seeds are clearly identified in the analysis.

The v5 frozen contract fixes sizes 128, 512 and 1,024; families clustered, independent and global; and seeds 55011–55013, with 128 training, 128 validation and 256 held-out epochs per workload. The executable fixture uses `SeedSequence([83012, seed, stream])`. Interrupted fixture runs append completed snapshot records to a journal and replay those records when resuming; they do not reclassify incomplete work as a completed seed.

The main commands, from the extracted research root, are:

```
python -m pytest -q
python -m experiments.synthetic --family clustered --n 128 --seeds 20
python -m experiments.executable --seeds 6 --out results/executable_fresh
python -m experiments.sensitivity
python -m experiments.aggregate
python -m experiments.validate_results
```

Nothing in these commands contacts an LLM endpoint. The artifact separates raw results from regenerated tables.

A.2 Software and data identities

Archived runs used Python 3.13.5 on Linux x86_64; exact installed package and platform records accompany emitted reports. Compilers were measured with `OPENBLAS_NUM_THREADS=1`. No accelerator is required. Vendored upstream sources and tests are retained with their notices. Every fresh mutation is checked against its archived base-source hash, replacement span, and resulting source hash. This binds a result to a specific intervention, not a natural historical bug distribution.

B Certificate Encoding and Trust-Boundary Details

Each capsule contains `range`, `counts`, `binding`, `evidence`, and `registry`. Hash fields are 64-character lowercase SHA-256 hexadecimal encodings. Canonical JSON uses sorted keys, ASCII escapes, and compact separators. In a packet, the header occupies 512 bytes including its newline; each child card occupies 320 bytes including its newline. Padding makes byte accounting exact. Fixed padding is not asserted to be a token-optimal production encoding; a binary or tokenizer-aware encoding would define a new price function and require re-evaluation.

The trusted verifier API rejects mismatched bindings, invalid verdicts, and malformed trace hashes. It is intentionally not exposed as a tool through which an untrusted LLM can certify its own work. Returned display capsules are compared with canonical trusted records rather than accepted because their self-reported hashes appear valid. This local design demonstrates record integrity and freshness under its trust assumptions; distributed authenticated transport and credential isolation remain host-runtime responsibilities.

The optimizer’s $O(bn^3)$ claim concerns planning. The reference kernel prioritizes clarity and scans interval generation stamps; it also hashes tree descriptions. We do not infer an asymptotically optimal runtime-update implementation from the communication bound. Large-scale implementations could use range-maximum indexes and cached topology digests, but those optimizations are outside the evaluated prototype.

C Optional Fixed-Share Analysis: Not the Deployment Recommendation

This appendix retains the former tracking theorem and its assumptions. It is an optional comparator, not the practical selector. On the measured streams it loses to the deterministic priced-window rule in three regimes. We have not searched for a favorable adversarial stream after seeing those results, and make no claim that a tracking bound implies an empirical advantage.

Let a declared bound R cover each round’s loss spread, $\max_i \ell_t(i) - \min_i \ell_t(i) \leq R$. Reject a measurement exceeding it instead of clipping the cost. Define $z_{t,i} = (\ell_t(i) - \min_j \ell_t(j))/R$. Starting at $p_1(i) = 1/K$, update

$$q_{t+1}(i) = \frac{p_t(i) \exp(-\eta z_{t,i})}{\sum_j p_t(j) \exp(-\eta z_{t,j})}, \quad (19)$$

$$p_{t+1}(i) = (1 - \alpha)q_{t+1}(i) + \alpha/K, \quad 0 < \eta \leq 1, \ 0 \leq \alpha < 1. \quad (20)$$

Independent draws from each new distribution can waste migrations. Instead, conditional on $A_t = i$, retain i with probability $\min\{1, p_{t+1}(i)/p_t(i)\}$. If it is not retained, sample from the normalized positive residual $(p_{t+1} - p_t)^+$. This maximal coupling has marginal p_{t+1} and switch probability $\text{TV}(p_t, p_{t+1})$. Established dependent-sampling methods motivate this choice [10]; we do not claim it as a new sampling algorithm.

Theorem 5 (Catalogue tracking with migration). *Assume a frozen catalogue, a fixed exogenous sequence of full-information loss vectors, spread at most R , and $M_{ij} \leq M$. For any comparator path $u_{1:T}$ with s changes and $0 < \alpha < 1$, the coupled policy satisfies*

$$\mathbb{E}[L_T] - \sum_t \ell_t(u_t) \leq R \left(\frac{\log K + s \log(K/\alpha) + (T-1-s) \log(1/(1-\alpha))}{\eta} + \frac{\eta T}{8} \right) + M(T-1)(\eta + \alpha). \quad (21)$$

For $\alpha = 0$ the stationary-comparator bound has numerator $\log K$ and applies only to $s = 0$.

Proof. The fixed-share update is exponential weighting over paths with an initial uniform prior, self-transition mass at least $1 - \alpha$, and transition mass at least α/K . The negative log prior of a path with s switches is bounded by the numerator. The bounded exponential-moment inequality gives expected normalized service regret at most that numerator divided by η , plus $\eta T/8$. Since losses are exogenous, the weights are independent of the sampled action path, and the coupling preserves each unconditional marginal. Exponential reweighting changes total variation by at most η ; uniform mixing adds at most α . Thus the expected switch count is at most $(T-1)(\eta + \alpha)$. Multiplication by M and addition of service regret proves the bound. $\square$ $\square$

The bound is not a per-run improvement guarantee. Choosing fixed $\eta = .5$ and $\alpha = .01$, as in our frozen experiment, need not make this bound sublinear with increasing T .

D Block-Backbone Compiler: Retained Ablation

The exact compiler’s cubic cost makes frequent full recompilation inappropriate. For a supplied display order, divide the leaves into consecutive blocks of at most s leaves, and fix a compact b -ary macro-tree on those blocks. Optimize each block’s microtree using its restricted completion waves and strict-epoch initialization cost. Replace each macro leaf by its optimized microtree. No evaluation waves enter this process.

The proposition below is exact only under its stated restrictions, and one restriction deserves emphasis because it is easy to misread as a general claim: *every update wave that intersects a macro node’s span must cover whole blocks within it*. If a wave can hit a partial block—some

but not all of the leaves a macro node covers—then that node’s activation weight varies within the class and the objective no longer separates into per-block subproblems; the $O(bns^2)$ claim then fails and the decomposition is only a heuristic. The frozen streams of this paper never emit partial-block waves (a wave’s gate set is a union of whole invalidation sets, and the measured waves intersect macro spans block-completely), which is why the decomposition is exact on them; a deployment with partial-block waves must either re-block or fall back to the full exact compiler.

Proposition 6 (Exactness in the block-backbone class). *With display order, block boundaries and macro-tree fixed, independent exact optimization of the blocks minimizes the training objective over every layout in that restricted class. The dynamic-programming work is $O(bns^2)$ with $O(bs^2)$ reusable DP workspace, excluding order learning and wave preparation.*

Proof. Every macro node covers a fixed union of complete blocks, so its activation weight and packet price are constant across the class. All variable costs lie inside disjoint microtrees. The objective is a constant plus the sum of their independent objectives. Conditional exactness therefore gives the minimizing member. There are at most $\lceil n/s \rceil$ subproblems, each of size at most s , yielding the stated DP bound. Reusing the subproblem arrays yields the workspace bound. $\square$ $\square$

For the reference implementation, M recorded waves and L total wave memberships add filtering and weight-preparation work of $O(\lceil n/s \rceil L + Mns)$, and storage for the input, output and local $M \times s$ arrays. This is not a linear-time claim for the entire pipeline. No bounded approximation ratio against the unrestricted exact tree is asserted; held-out cost can deteriorate. The historical experiment froze $s = 32$ before measurement. Its negative result—block DP lost to same-order balanced fixed8 in held-out bytes—motivates excluding this compiler from the practical path.

E Information Separations

These results concern the information available to a selector. They are not hardness results for all adaptive tree algorithms, and they do not imply that a complete hypergraph must be recovered: interval hit statistics are sufficient for a fixed order and objective.

E.1 Marginals do not determine a co-invalidation optimum

Consider four ordered leaves, binary trees, and a constant internal packet price $c > 0$. Let $\mathcal{D}_A$ choose $\{0, 1\}$ or $\{2, 3\}$ equiprobably and $\mathcal{D}_B$ choose $\{0, 3\}$ or $\{1, 2\}$ equiprobably. All leaf marginal activation probabilities equal $1/2$ in both distributions. The root and every three-leaf interval are always activated. The five tree shapes fall into the following classes:

Ordered tree	J_A/c	J_B/c
$((0, 1), (2, 3))$	2	3
$(0, ((1, 2), 3)), ((0, (1, 2)), 3)$	3	2.5
$(0, (1, (2, 3))), (((0, 1), 2), 3)$	2.5	3

Optimal costs are $2c$ and $2.5c$. A deterministic marginal-only selector receives identical input and must choose the same class, incurring at least $c/2$ regret on one distribution. The last class is weakly dominated by the balanced tree. If a randomized selector chooses the balanced tree with probability q , its regrets are at least $(1 - q)c$ and $qc/2$. Their maximum is minimized at $q = 2/3$ and equals $c/3$, which is attainable for this instance. A fixed initialization price adds the same $3c$ to all binary shapes and does not alter the separation.

E.2 Joint invalidation does not determine progressive layout

Every epoch now invalidates all four checks, and execution order is $(0, 1, 2, 3)$ with singleton completion waves. In workload A , check 0 fails, so only its wave is executed. In workload B , check 3 fails, so all four checks execute. Initial invalidation is identical, even jointly. Every binary tree has three internal nodes; its cost in A is $(3 + \text{depth}(0))c$ and in B is $(3 + \sum_i \text{depth}(i))c$. The root-0 comb minimizes A at $4c$ but costs $12c$ in B , whereas the balanced tree minimizes B at $11c$ but costs $5c$ in A . An invalidation-only randomized selector using a root-0 comb with probability p incurs regrets at least $(1 - p)c$ and pc , whose maximum is at least $c/2$. A deterministic selector incurs at least c . This proves a separation between stale-set information and schedule-induced completion information; it does not prohibit a selector that observes both.

F Cost, Statistical, and Online Proof Details

F.1 Full-refresh lower bound and global negative control

For an n -leaf tree with m internal nodes, the number of edges is $n + m - 1$, so $\sum_v d_v = n + m - 1$. Therefore

$$W(L) = S(n - 1) + (H + S)m \geq S(n - 1) + (H + S) \left\lceil \frac{n - 1}{b - 1} \right\rceil. \quad (22)$$

The inequality follows from $n - 1 = \sum_v (d_v - 1) \leq m(b - 1)$. At $n = 64$, $b = 8$, a complete eight-child construction attains the bound with $m = 9$ and $W = 27,648$ bytes. Thus all selected global-control layouts can tie an analytically minimal full refresh. Strict initialization followed by one global completion wave costs $2W$. This explains the exact negative control without assuming every balanced tree at every n attains the lower bound.

F.2 Uniform statistical control at a fixed order

For one activation wave per episode and a fixed registry order, let $M_n = n(n + 1)/2$ and $\epsilon = \sqrt{\log(2M_n/\delta)/(2m)}$. The interval hit indicators are bounded independent variables across independent episodes, though checks within an episode may be arbitrarily dependent. Hoeffding plus a union bound gives simultaneous interval-frequency error at most ϵ . For every feasible tree, $|\hat{J}(T) - J(T)| \leq W(T)\epsilon$. Inserting the empirical minimizer gives regret at most $[W(\hat{T}) + W(T^*)]\epsilon \leq 2(n - 1)C\epsilon$. This is the original HACT bound, retained for attribution. It is not automatically valid uniformly over arbitrary data-selected permutations; the separate held-out catalogue bound in Proposition 3 addresses that selection step.

For strict epochs with at most q completion waves, a loose range is $B = (q + 1)(n - 1)C$. Under the evaluated sequential eight-record flush rule with at most one final record per obligation, $q \leq \lceil n/8 \rceil$. A narrower observed range is not automatically a proved bound for an unobserved deployment distribution.

F.3 Maximal coupling and horizon dependence

For consecutive distributions p, p' , the probability mass that remains at index i under the coupling is $\min(p_i, p'_i)$. The remaining mass is distributed as $(p' - p)^+$, so the new marginal is exactly p' and $\Pr(A_{t+1} \neq A_t) = \text{TV}(p, p')$. For exponential reweighting by $a_i = \exp(-\eta z_i)$ with $z_i \in [0, 1]$, the normalizer is at most one, hence $q_i \geq e^{-\eta} p_i$. Thus $\sum_i \min(p_i, q_i) \geq e^{-\eta}$ and $\text{TV}(p, q) \leq 1 - e^{-\eta} \leq \eta$. Mixing q with the uniform distribution by weight α changes total variation by at most α . The triangle inequality gives the switch bound in Theorem 5.

For $K \geq 2$, $s = 0$, and $\alpha = 0$, write $B_T = RT/8 + M(T - 1)$. The bound is $R \log K/\eta + \eta B_T$. When $\sqrt{R \log K/B_T} \leq 1$, choosing that learning rate gives at most $2\sqrt{R \log K B_T}$ expected regret,

sublinear for fixed R, M, K . This is an established style of switching-cost bound, not a new rate. A single empirical switch with estimated future saving d and anticipated horizon h satisfies the familiar break-even inequality $hd > M_{ij}$. This is arithmetic, not a drift theorem: savings may reverse before h events.

G Evidence-to-Claim Register

Claim	Evidence	Scope boundary
Exact ordered optimum	Theorem 2; exhaustive enumeration tests	Fixed order and price model
Validation after order	Proposition 3; Hoeffding bound	Fixed catalogue, no shift
Publication soundness	Theorem 4; conditional proof	Trusted checker, protected state
Practical byte reduction	Table 6; frozen seeds	Clustered workloads, 3 seeds
Source replay reductions	Table 7; 36 replays	Diagnostic, not fresh checks
Root status identical	Equation (15); 593–594 bytes	Fixed tokenizer, layout-independent
Break-even accounting	Equations (17)–(18)	Declared C, E, R , codec-matched
BPE tokenization gap	c1100k_base/o200k_base remain UNKNOWN; admission tool supplied	Vocabulary retrieval failed; no billing claim
Context-retirement mechanism study	Pre-registered, factorial, executed: v2 = 1,176 runs + 63 long-horizon, seven retention levels × scheduler factor, five closures (two proposals rejected by their gates)	Measured at mechanism level; live cohort still protocol-only

H Retirement-Study Protocol Deviations and Engine History

Two protocol generations were run. Generation v1 (file-frozen 2026-09-23; docs/protocols/CONTEXT_RETIREMENT_PROTOCOL_v1.md with amendments v1.1–v1.4) was written before cohort data generation, and each amendment was written before the run it governs. The engine went through four substantive revisions during that pre-data window; the final cohort was generated once, by the amended engine, and never regenerated after any manuscript text was written. For auditability we disclose the full history, including the errors that motivated each amendment:

- Amendment v1.1 (pre-data).** The HACT workload generators draw invalidation sets i.i.d.; retirement is temporal, so persistence transitions (0.6) and fixed obligation windows (six episodes) were added. A single-stream smoke test of the original design exposed nominal (name-based) obligation retention; the amendment made retention emergent from each policy’s mechanism.
- Amendment v1.2 (pre-data).** A 630-run smoke test to a temporary directory (results void, never cited) exposed three under-specifications: hand-rolled evidence accounting, summary/reset modeled as sleeping between resets rather than attending, and a frontier capsule that did not scale with contents. All were pinned before the final run.
- Amendment v1.3 (pre-data).** A single-stream check showed facts carrying no payload, making replay artificially cheap; payloads (lognormal, median 128 B), payload-charged retrieval, and interface-only capsules were pinned.

4. **Amendment v1.4 (pre-data).** Final verification exposed missing dependency footprints (the review’s exported interfaces), an inconsistent summary/reset failure model, and the need for retrieval operations as a co-primary outcome. The evidence ablation was refined to a held-out half with a learned layout order and a steel-manned flat ledger taking the better of full and delta exports.
5. **Post-freeze fixes.** A flat-ledger registry-size off-by-one was corrected after the first final-format cohort; the cohort was regenerated once, before any manuscript table was generated from it. One obligation-resolution accounting error (retained-value reads not crediting resolution) was fixed during pre-data engine verification; all obligation flows are now covered by the test suite (`tests/retirement`), which also pins determinism, the closure outcomes, and the per-family evidence directions.

No amendment was written after the final cohort was generated. The v1 cohort, closures, and summary JSON are reproduced by `bun src/harness/retirement/cli.ts` in under two minutes and are byte-identical across runs; the v2 cohort, closures, and tables are reproduced by `bun src/harness/retirement/run2.ts` and `bun src/harness/retirement/tex2.ts` with the same property. Both engines, both protocols, and both frozen cohorts are committed; v2’s protocol and engine-notes commits precede its data commit, which precedes the first v2 manuscript-text commit.

I Generation-v1 Record (Complete)

The v1 study is superseded by the v2 factorial (Section 4) but is retained in full for provenance: v2’s design was chosen to repair v1’s weaknesses, and v1’s closures motivated v2’s frozen rules. The v1 protocol was frozen in files on 2026-09-23, before its cohort ran; it was first committed to git on 2026-09-24, after the v2 protocol commit (333d495). v1’s freeze is therefore file-attested, and the ordering is disclosed rather than smoothed.

Design. Ten arms over $3 \times 3 \times 7 = 63$ cells (630 runs, 40 episodes each): full-replay, masking, summary-reset, bounded-hier, verified-board, frontier, frontier+adapt, adapt-only, frontier+flat, frontier+hact. Three closure rules: null (frontier $\leq 80\%$ of best-baseline cost-per-accepted), attribution (frontier+adapt within 10% of adapt-only), HACT (tree $\leq 95\%$ of flat steel-man on held-out evidence bytes).

Frozen outcomes (as generated 2026-09-23, never recomputed).

- *Null closure triggered:* frontier cost-per-accepted 1,906,765 versus best baseline 564,773 (threshold $\leq 451,818$). No byte-cost claim for frontier retirement.
- *Summary/reset falsified:* 21/63 accepted, 30,820 stale reads, 280 lost obligations—the lossy-conversion failure mode.
- *Operations:* frontier cut retrieval operations 48.3% versus masking (27,728 vs 53,583 across the cohort), the observation that became v2’s regime-O hypothesis.
- *Attribution closure triggered:* scheduler dimension byte-neutral (frontier+adapt identical to frontier on bytes); read dimension-separated (makespan only; 87.4% fewer rounds with adaptive batching).
- *HACT closure not triggered:* held-out evidence bytes 59.8% lower than the flat steel-man pooled; per-family split (+88.7% clustered, +49.2% independent, −87.1% global) forced the conditional reading formalized in v2’s density rule.

Why v2 replaced the design. v1’s arms were hand-picked combinations, so factor attribution leaned on pairwise comparisons and a single attribution closure; its null rule read one pooled byte threshold; its capsule accounting was an end-of-run total with no independent audit; and it priced one cost regime, which is why the operations result could only be reported as a hypothesis. Each weakness became a v2 design item (Section 4.2). v1’s cohort, closures, and summary remain reproducible byte-identically from the committed v1 engine via `bun src/harness/retirement/cli.ts`.

References

- [1] Gregg Roethermel, Roland H. Untch, Chengyun Chu, and Mary Jean Harrold. Prioritizing test cases for regression testing. *IEEE Transactions on Software Engineering*, 27(10):929–948, 2001. <https://doi.org/10.1109/32.962562>.
- [2] Milos Gligoric, Lamya Eloussi, and Darko Marinov. Practical regression test selection with dynamic file dependencies. *Proceedings of ISSSTA*, 2015.
- [3] You Wang, Michael Pradel, and Zhongxin Liu. Names Are All You Need: Effective and Safe Regression Test Selection for Python. arXiv:2605.25356v1, May 2026.
- [4] Umut Acar, Guy Blelloch, Matthias Blume, Robert Harper, and Kanat Tangwongsan. A Library for Self-Adjusting Computation. *Electronic Notes in Theoretical Computer Science*, 148:127–154, 2006.
- [5] Andrey Mokhov, Neil Mitchell, and Simon Peyton Jones. Build Systems à la Carte. *Proceedings of the ACM on Programming Languages*, 2(ICFP), Article 79, 2018.
- [6] Hirendu Vaishnav and Massoud Pedram. Alphabetic Trees: Enumeration and Optimization with Applications to VLSI Design. Author-hosted manuscript, Sections 4.1–4.4; publication date not established from the retrieved copy.
- [7] Ziheng Shangguan, Aviv Yaish, and Dahlia Malkhi. Authenticated Data Structures for Dynamic Workloads. arXiv:2608.25206v1, August 2026.
- [8] Jonathan E. Atkins, Erik G. Boman, and Bruce Hendrickson. A Spectral Algorithm for Seriation and the Consecutive Ones Problem. *SIAM Journal on Computing*, 28(1):297–310, 1998.
- [9] Mark Herbster and Manfred K. Warmuth. Tracking the Best Expert. *Machine Learning*, 32:151–178, 1998.
- [10] Sascha Geulen, Berthold Vöcking, and Melanie Winkler. Regret Minimization for Online Buffering Problems Using the Weighted Majority Algorithm. *Proceedings of COLT*, pages 132–143, 2010.
- [11] Nicolò Cesa-Bianchi, Ofer Dekel, and Ohad Shamir. Online Learning with Switching Costs and Other Adaptive Adversaries. *Advances in Neural Information Processing Systems*, 26, 2013.
- [12] Zexun Wang. Proof-Carrying Agent Actions: Model-Agnostic Runtime Governance for Heterogeneous Agent Systems. arXiv:2606.04104v1, June 2026.
- [13] Zihao Zheng, Baichuan Li, Junyi Yao, and Jiayu Long. Engineering Reliable Commit Gates for Agentic AI: Cost-Aware Verification Portfolios under Common-Mode Data Failures. arXiv:2609.10969v1, September 2026.

- [14] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? *International Conference on Learning Representations*, 2024.
- [15] OpenTelemetry contributors. Batch Processor documentation, official collector implementation. <https://github.com/open-telemetry/opentelemetry-collector/tree/main/processor/batchprocessor>.
- [16] Lindenbauer et al. The Complexity Trap: Simple Observation Masking Is as Efficient as LLM Summarization for Agent Context Management. arXiv:2508.21433v3, October 2025.
- [17] Sun et al. Scaling Long-Horizon LLM Agent via Context-Folding. arXiv:2510.11967.
- [18] Chen et al. Towards a Science of Scaling Agent Systems (v3). Companion review reference [R4]; controlled scaling results depend on task-architecture alignment.
- [19] Yuzhen Mao and Azalia Mirhoseini. Decentralized Multi-Agent Systems with Shared Context. arXiv:2606.10662v1, June 2026.
- [20] OpenAI contributors. tiktoken: BPE tokenization implementation and encoding registry. <https://github.com/openai/tiktoken>.
- [21] SWE-bench contributors. The Harness: Docker-based patch evaluation reference. <https://www.swebench.com/SWE-bench/reference/harness/>.
- [22] B. Constantine, G. Forget, R. Geib, and R. Schrage. Framework for TCP Throughput Testing. RFC 6349, August 2011.
- [23] Wassily Hoeffding. Probability Inequalities for Sums of Bounded Random Variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.